

Regression analysis with scikit-learn

HAS Tools - Sept 29

DISCLAIMER


We are playing fast and loose with some stats today, please be sure to follow up on formal methods either by taking some real stats classes or using prior knowledge!!!

Where we left off on Tuesday

You saw how to load data from the USGS streamflow database for the Verde River

You also learned how to apply the same technique to the DayMet dataset, which lets us get things like temperature and precipitation

We finished up by making this plot




Where we left off on Tuesday

You saw how to load data from the USGS streamflow database for the Verde River

You also learned how to apply the same technique to the DayMet dataset, which lets us get things like temperature and precipitation

We finished up by making this plot

Clearly there's some relationship here, but can we quantify it?




Where we left off on Tuesday

You saw how to load data from the USGS streamflow database for the Verde River

You also learned how to apply the same technique to the DayMet dataset, which lets us get things like temperature and precipitation

We finished up by making this plot

Clearly there's some relationship here, but can we quantify it?



Enter linear regression


We're going to do just a little bit of theory here before

Suppose we have data with:

- $x = [1, 2, 3, 4]$
- $y = [6, 5, 7, 10]$

We will call x the “explanatory variable”

And we call y the “response variable”



So how do we go about fitting a line through the data?


Recall that the equation for a line is:

$$y = \beta_2 x + \beta_1$$

Where:

β_2 Is the slope

β_1 Is the intercept



So how do we go about fitting a line through the data?

Recall that the equation for a line is:


$$y = \beta_2 x + \beta_1$$

Where:

β_2 Is the slope

β_1 Is the intercept

r Is a residual term




Now we can frame things into a system of equations

$$\beta_1 + 1\beta_2 + r_1 = 6$$

$$\beta_1 + 2\beta_2 + r_2 = 5$$

$$\beta_1 + 3\beta_2 + r_3 = 7$$

$$\beta_1 + 4\beta_2 + r_4 = 10$$



Now we can frame things into a system of equations


We want the value of these

$$\beta_1 + 1\beta_2 + r_1 = 6$$

$$\beta_1 + 2\beta_2 + r_2 = 5$$

$$\beta_1 + 3\beta_2 + r_3 = 7$$

$$\beta_1 + 4\beta_2 + r_4 = 10$$




Now we can frame things into a system of equations

We want the value of these

$$\begin{array}{rclcl} \beta_1 + 1\beta_2 + r_1 & = & 6 \\ \beta_1 + 2\beta_2 + r_2 & = & 5 \\ \beta_1 + 3\beta_2 + r_3 & = & 7 \\ \beta_1 + 4\beta_2 + r_4 & = & 10 \end{array}$$

And we want to minimize the
“total” of these




We rearrange to put the residual terms on the LHS

$$r_1 = 6 - (\beta_1 + 1\beta_2)$$

$$r_2 = 5 - (\beta_1 + 2\beta_2)$$


$$r_3 = 7 - (\beta_1 + 3\beta_2)$$

$$r_4 = 10 - (\beta_1 + 4\beta_2)$$



Then we reframe in terms of a sum of squares equation


$$\begin{aligned} S(\beta_1, \beta_2) &= r_1^2 + r_2^2 + r_3^2 + r_4^2 \\ &= [6 - (\beta_1 + 1\beta_2)]^2 \\ &\quad + [5 - (\beta_1 + 2\beta_2)]^2 \\ &\quad + [7 - (\beta_1 + 3\beta_2)]^2 \\ &\quad + [10 - (\beta_1 + 4\beta_2)]^2 \end{aligned}$$



Then we reframe in terms of a sum of squares equation

$$\begin{aligned} S(\beta_1, \beta_2) &= r_1^2 + r_2^2 + r_3^2 + r_4^2 \\ &= [6 - (\beta_1 + 1\beta_2)]^2 \\ &\quad + [5 - (\beta_1 + 2\beta_2)]^2 \\ &\quad + [7 - (\beta_1 + 3\beta_2)]^2 \\ &\quad + [10 - (\beta_1 + 4\beta_2)]^2 \end{aligned}$$


$$= 4\beta_1^2 + 30\beta_2^2 + 20\beta_1\beta_2 - 56\beta_1 - 154\beta_2 + 210$$



Set the partial derivatives equal to zero to find the “least squares error”

$$\frac{\partial S}{\partial \beta_1} = 0 = 8\beta_1 + 20\beta_2 - 56$$

$$\frac{\partial S}{\partial \beta_2} = 0 = 20\beta_1 + 60\beta_2 - 154.$$




Set the partial derivatives equal to zero to find the
“least squares error”

$$\frac{\partial S}{\partial \beta_1} = 0 = 8\beta_1 + 20\beta_2 - 56$$

$$\frac{\partial S}{\partial \beta_2} = 0 = 20\beta_1 + 60\beta_2 - 154.$$

***2 equations, 2 unknowns,
easy-peasy***




Set the partial derivatives equal to zero to find the
“least squares error”

$$\beta_1 = 3.5$$

$$\beta_2 = 1.4$$

***2 equations, 2 unknowns,
easy-peasy***




Set the partial derivatives equal to zero to find the
“least squares error”

$$\beta_1 = 3.5$$

$$\beta_2 = 1.4$$

***2 equations, 2 unknowns,
easy-peasy***




Looks pretty good! But I'm glossing over **A TON**

The statistical procedure for fitting Ordinary Least Squares (OLS) models is generally understood

However there are many caveats, conditions, etc

And, particularly, linear least squares is not always the best model (actually almost never)




Looks pretty good! But I'm glossing over **A TON**

The statistical procedure for fitting Ordinary Least Squares (OLS) models is generally understood

However there are many caveats, conditions, etc

And, particularly, linear least squares is not always the best model (actually almost never)



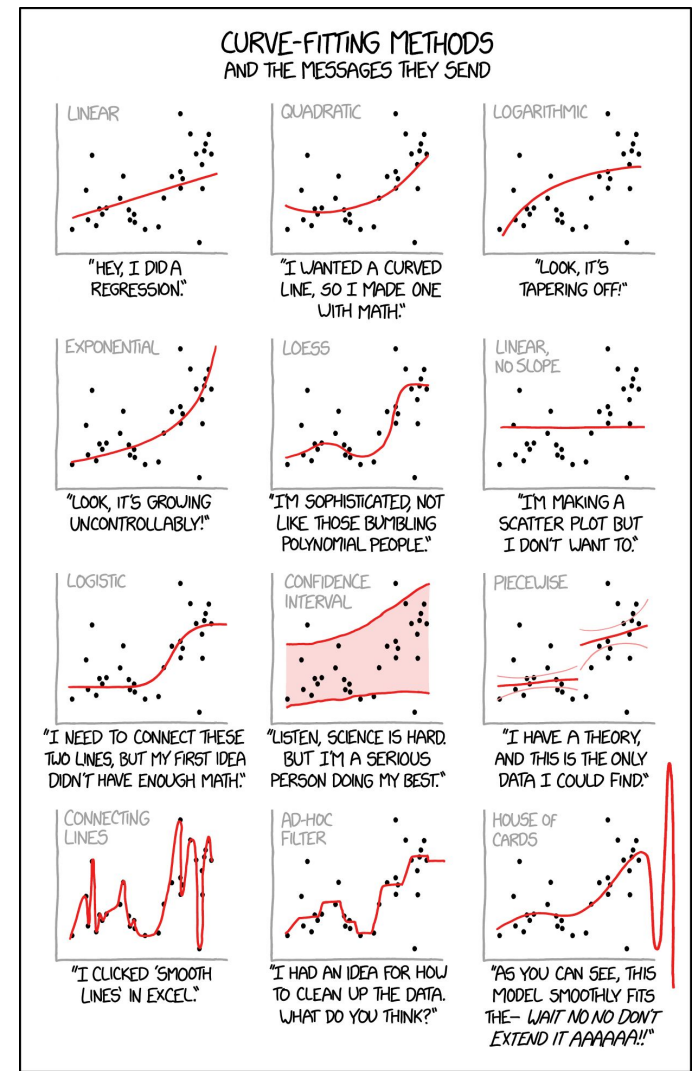
Looks pretty good! But I'm glossing over **A TON**

The statistical procedure for fitting Ordinary Least Squares (OLS) models is generally understood

However there are many caveats, conditions, etc

And, particularly, linear least squares is not always the best model (actually almost never)

And then there's non-linear regression....





**HOW DO WE
DO THIS WITH
PYTHON!?**

The big 3



The big 3



Why scikit-learn?



It is well established, same for statsmodels & scipy

It has a consistent interface (or API) between model types

It has a number of other applications I hope you find useful




Classification

Identifying which category an object belongs to.

Applications: Spam detection, Image recognition.

Algorithms: SVM, nearest neighbors, random forest, and more....




Examples

Regression

Predicting a continuous-valued attribute associated with an object.

Applications: Drug response, Stock prices.

Algorithms: SVR, nearest neighbors, random forest, and more....



Examples

Clustering

Automatic grouping of similar objects into sets.

Applications: Customer segmentation, Grouping experiment outcomes

Algorithms: k-Means, spectral clustering, mean-shift, and more....




Examples

Dimensionality reduction

Reducing the number of random variables to consider.

Applications: Visualization, Increased efficiency

Algorithms: PCA, feature selection, non-negative matrix factorization, and more....




Examples

Model selection

Comparing, validating and choosing parameters and models.

Applications: Improved accuracy via parameter tuning

Algorithms: grid search, cross validation, metrics, and more....




Examples

Preprocessing

Feature extraction and normalization.

Applications: Transforming input data such as text for use with machine learning algorithms.

Algorithms: preprocessing, feature extraction, and more....



Examples




Classification

Identifying which category an object belongs to.

Applications: Spam detection, Image recognition.

Algorithms: SVM, nearest neighbors, random forest, and more....




Examples

Regression

Predicting a continuous-valued attribute associated with an object.

Applications: Drug response, Stock prices.

Algorithms: SVR, nearest neighbors, random forest, and more....



Examples

Clustering

Automatic grouping of similar objects into sets.

Applications: Customer segmentation, Grouping experiment outcomes

Algorithms: k-Means, spectral clustering, mean-shift, and more....




Examples

Dimensionality reduction

Reducing the number of random variables to consider.

Applications: Visualization, Increased efficiency

Algorithms: PCA, feature selection, non-negative matrix factorization, and more....




Examples

Model selection

Comparing, validating and choosing parameters and models.

Applications: Improved accuracy via parameter tuning

Algorithms: grid search, cross validation, metrics, and more....




Examples

Preprocessing

Feature extraction and normalization.

Applications: Transforming input data such as text for use with machine learning algorithms.

Algorithms: preprocessing, feature extraction, and more....



Examples



Jump to VSCode,
time to install
scikit-learn into conda

https://scikit-learn.org/stable/auto_examples/linear_model/plot_ols.html